

# Profiling

CMSC 257: Computer Systems

**Instructor:**

Preetam Ghosh ([pghosh@vcu.edu](mailto:pghosh@vcu.edu))

# Program Performance

- Programs run only as well as the code you write
- Poor code often runs poorly
  - Crashes or generates incorrect output (bugs)
  - Is laggy, jittery or slow (inefficient code)
    - Too slow on real inputs (data processing)
    - Not-reactive enough to be usable (interfaces)

# Optimization

- Optimization is the process where you take an existing program and alter it to remove inefficiencies.
  - Change algorithms
  - Restructure code
  - Redesign data structures
  - Refactor code

# Example Inefficient Program

```
/* Program: profiling */
#include <stdio.h>

int add_example_one(int a, int b) {
    int out, i, j;
    out = 0;
    for (i=0; i<a; i++) {
        out ++;
    }
    for (j=0; j<b; j++) {
        out ++;
    }
    return(out);
}

int add_example_two(int a, int b) {
    return( a+b );
}

int main(void) {
    int i, x, y;
    for (i=0; i<10000; i++) {
        x = add_example_one(10000, 20000);
        y = add_example_two(10000, 20000);
    }
    return(0);
}
```

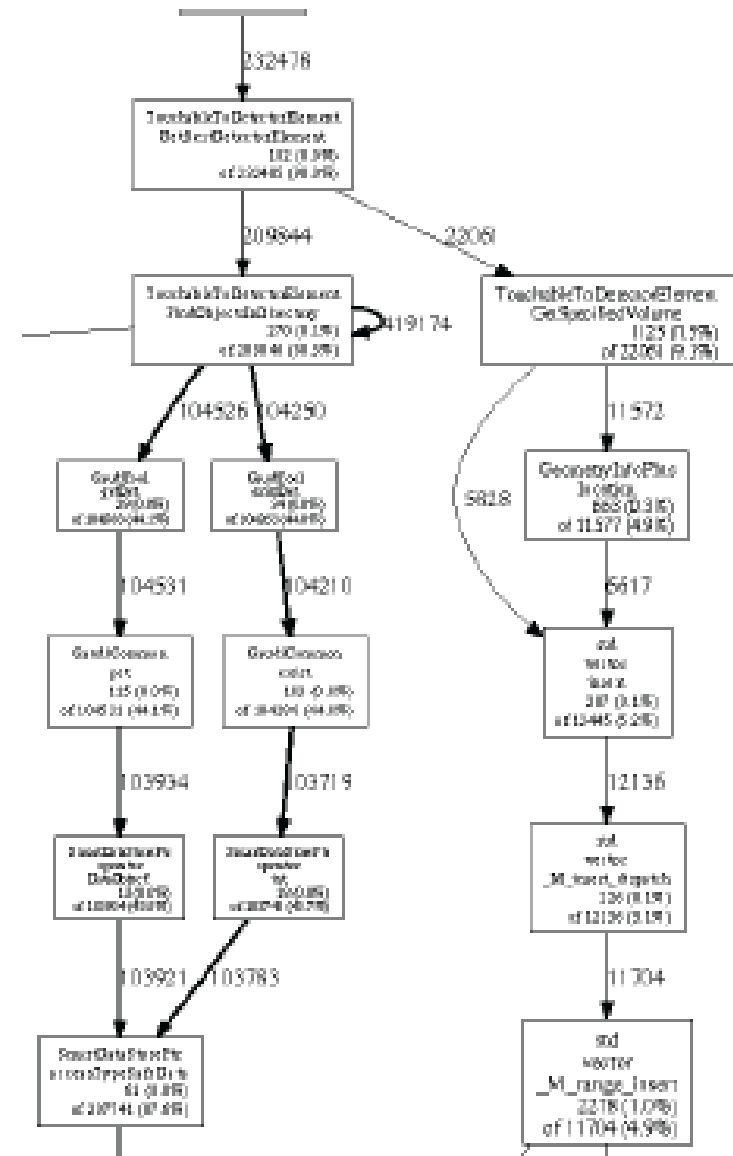
- Q: Which function is going to run faster and why?

# Profiling

- Debugging helps the programmer find and fix bugs ...
- *Profiling* helps the programmer find inefficiencies
  - Profiling involves running a version of the program instrumented with code to measure how much time is spent in certain areas of the code.
  - How much time in each of the modules of the program?

# gprof

- **gprof** is a utility that measures a program's performance and behavior.
- This produces *non-interactive* profile output
- The output provides detail
  - The time that the program ran, and time for each function
    - Statistics and detail on “*performance*”
  - Which functions called other functions and how many times
    - Statistics and detail on the “*call graph*”



# Running gprof

1. First compile program using the “-pg” flag:  
`$ gcc -pg profiling.c -o profiling`
2. Run the program (will generate file gmon.out):  
`$ ./profiling`
3. Run gprof with the named program  
`$ gprof profiling | less`
4. Review the output

A terminal window with a black background and yellow and blue text. The text shows the command '\$ gprof profiling' followed by the output 'Flat profile:' and three dots '...'.

```
$ gprof profiling
Flat profile:
...
```

5. Optimize the program, re-profile
6. GOTO step#1

# Gprof (flat profile)

```
$ gprof profiling
```

```
Flat profile:
```

```
Each sample counts as 0.01 seconds.
```

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
101.15	0.47	0.47	10000	46.53	46.53	add_example_one
0.00	0.47	0.00	10000	0.00	0.00	add_example_two

%  
time      the percentage of the total running time of the  
program used by this function.

cumulative  
seconds    a running sum of the number of seconds accounted  
for by this function and those listed above it.

self  
seconds    the number of seconds accounted for by this  
function alone. This is the major sort for this  
listing.

calls      the number of times this function was invoked, if  
this function is profiled, else blank.

self  
ms/call    the average number of milliseconds spent in this  
function per call, if this function is profiled,  
else blank.

total  
ms/call    the average number of milliseconds spent in this  
function and its descendents per call, if this  
function is profiled, else blank.

```
...
```



# Gprof (Call Graph)

```
...
Call graph (explanation follows)

index % time      self  children  called      name
-----
[1]    100.0      0.47    0.00    10000/10000    main [2]
      100.0      0.47    0.00    10000          add_example_one [1]
-----
                                <spontaneous>
[2]    100.0      0.00    0.47          main [2]
      0.47      0.00    10000/10000    add_example_one [1]
      0.00      0.00    10000/10000    add_example_two [3]
-----
[3]      0.0      0.00    0.00    10000/10000    main [2]
      0.00      0.00    10000          add_example_two [3]
-----
```

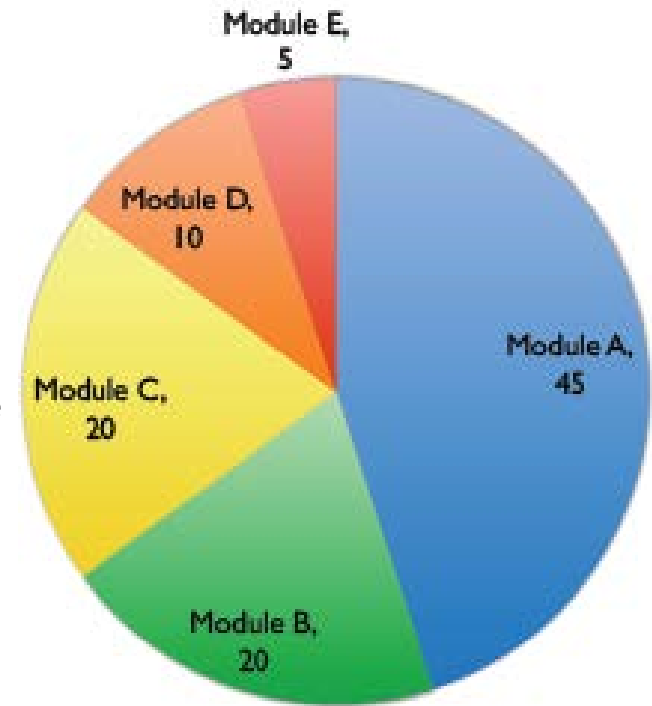
This table describes the call tree of the program, and was sorted by the total amount of time spent in each function and its children.

Each line lists:

% time	This is the percentage of the `total' time that was spent in this function and its children.
self	This is the total amount of time spent in this function.
children	This is the total amount of time propagated into this function by its children.
called	This is the number of times the function was called. If the function called itself recursively, the number only includes non-recursive calls, and is followed by a `+' and the number of recursive calls.

# Optimization revisited ...

- When optimizing, you focus on modules of the program which implement the features and processing of the program.
  - Which parts of the program you select depends on what parts are running the most.
    - then, focus on those parts which take up the most time.
- Profiling tells us where to spend our time.



# Amdahl's Law

- Amdahl's law models the maximum performance gain that can be expected by improving part of the system, i.e., what can we expect in terms of improvement.
- Consider
  - ▶ **k** is the percentage of total execution time spent in the optimized module(s).
  - ▶ **s** is the execution time expressed in terms of a n-factor speedup (2X, 3X...), which can be found as

$$\frac{\text{original execution time}}{\text{improved execution time}}$$

# Amdahl's Law (cont.)

- The overall speedup  $T$  of the program is expressed :

$$T = \frac{1}{(1 - k) + \frac{k}{s}}$$

- Intuition:

- ▶  $1 - k$  is the part of the program that is unchanged
- ▶  $\frac{k}{s}$  is the ratio of altered program size to speedup

# Amdahl's Law (example)

- Assume that a module A of a program is optimized.
  - A represents 45% of the run time of the program.
  - The optimization reduces the runtime of module from 750ms to 50ms.
  
- What is the program speedup?

# Amdahl's Law (example)

- Assume that a module A of a program is optimized.
  - A represents 45% of the run time of the program.
  - The optimization reduces the runtime of module from 750ms to 50ms.

$$s = 750/50 = 15$$

$$k = .45$$

$$T = \frac{1}{(1-.45) + \frac{.45}{15}} = \frac{1}{.55 + .03} = 1.724$$

- What is the program speedup? (A: **1.724X**)